

## NLP-LORA-FINETUNING AND BENCHMARK RESULTS

### 1. Project Overview

The objective of this project is to enhance the code reasoning and generation capabilities of a pre-trained Large Language Model (LLM) using Low-Rank Adaptation (LoRA). By fine-tuning the base model on two distinct datasets, I aim to compare the effects of “Deep Reasoning” traces versus “Diverse Problem Types” on the model’s final performance in competitive programming tasks.

- Base Model: Qwen/Qwen2.5-Coder-1.5B-Instruct
- Technique: Parameter-Efficient Fine-Tuning (PEFT) via LoRA.

### 2. Dataset Analysis

Two separate datasets were utilized for comparative training:

1. DEEP Dataset (CodeGen-Deep-5K): Focused on complex reasoning steps and logical consistency.
2. DIVERSE Dataset (CodeGen-Diverse-5K): Focused on a wide variety of coding problems and edge cases.

**Note:** For both models, training was focused on the solution field (clean Python code) to optimize for direct output quality.

### 3. Training Setup & Hyperparameters

The following technical configurations were applied to ensure efficient training while preventing overfitting.

#### 3.1. LoRA Configuration

##### Deep Dataset

Parameter	Value	Description
Rank (r)	16	Rank of the update matrices
Alpha	32	LoRA scaling factor
Target Modules	q_proj, v_proj, k_proj, o_proj, gate_proj, up_proj, down_proj	Layers where LoRA is applied

<b>Dropout</b>	0.1	Prevention of co-adaptation of weights
<b>Bias</b>	None	No bias terms were trained

### Diverse Dataset

Parameter	Value	Description
<b>Rank (r)</b>	32	Rank of the update matrices
<b>Alpha</b>	64	LoRA scaling factor
<b>Target Modules</b>	q_proj, v_proj, k_proj, o_proj, gate_proj, up_proj, down_proj	Layers where LoRA is applied
<b>Dropout</b>	0.05	Prevention of co-adaptation of weights
<b>Bias</b>	None	No bias terms were trained

### 3.2 Technical Configuration

#### Deep Dataset

<b>Optimizer</b>	AdamW (8-bit)
<b>Learning Rate</b>	2e-5 (with Cosine Decay)
<b>Epochs</b>	3
<b>Context Length</b>	1024
<b>Effective Batch Size</b>	16
<b>System Prompt</b>	"You are an expert Python programmer. Please read the problem carefully before writing any Python code."


## Diverse Dataset

<b>Optimizer</b>	AdamW (8-bit)
<b>Learning Rate</b>	1e-4 (with Cosine Decay)
<b>Epochs</b>	3
<b>Context Length</b>	1024
<b>Effective Batch Size</b>	16
<b>System Prompt</b>	"You are an expert Python programmer. Please read the problem carefully before writing any Python code."

## 4. Training Results

### 4.1. Loss Curves

The following charts illustrate the convergence of the training and validation loss over the steps.



Deep Model Loss Graph



Diverse Model Loss Graph

### 4.2. Checkpoint Selection

Checkpoints were evaluated every 100 steps. The best-performing checkpoints selected for final evaluation were:

- DEEP Model: checkpoint-step-846-epoch-3
- DIVERSE Model: checkpoint-step-800-epoch-3

## 5. Final Evaluation & Benchmarks

The models were evaluated against a benchmark dataset to measure their accuracy in generating valid, functional Python code.

Model	Best Checkpoint	Pass@1 (%)	Solved Questions
<b>Base Model (Qwen2.5-Coder-1.5B)</b>		21.9	9/41
<b>Fine-Tuned (DEEP Model)</b>	checkpoint-step-846-epoch-3	[Value]	10/41
<b>Fine-Tuned (DIVERSE Model)</b>	checkpoint-step-800-epoch-3	26.8	11/41

## 6. Conclusion

In this project, I successfully implemented LoRA fine-tuning on the Qwen2.5-Coder-1.5B-Instruct model. My finding suggests that ... This experimentation proves that parameter-efficient fine-tuning can significantly specialize a small-scale model for high-level technical tasks.

## 7. Resources & Links

- GitHub Repository: <https://github.com/abdurrahman-gulmez/NLP-LLM-FINETUNING>
- HuggingFace Dataset (DEEP): <https://huggingface.co/datasets/Naholav/CodeGen-Deep-5K>
- HuggingFace Dataset (DIVERSE): <https://huggingface.co/datasets/Naholav/CodeGen-Diverse-5K>
- HuggingFace Model (Qwen): <https://huggingface.co/Qwen/Qwen2.5-Coder-1.5B-Instruct>